

Class 7: Machine Learning 1

Ivan Henry Kish (PID:A17262923)

Background

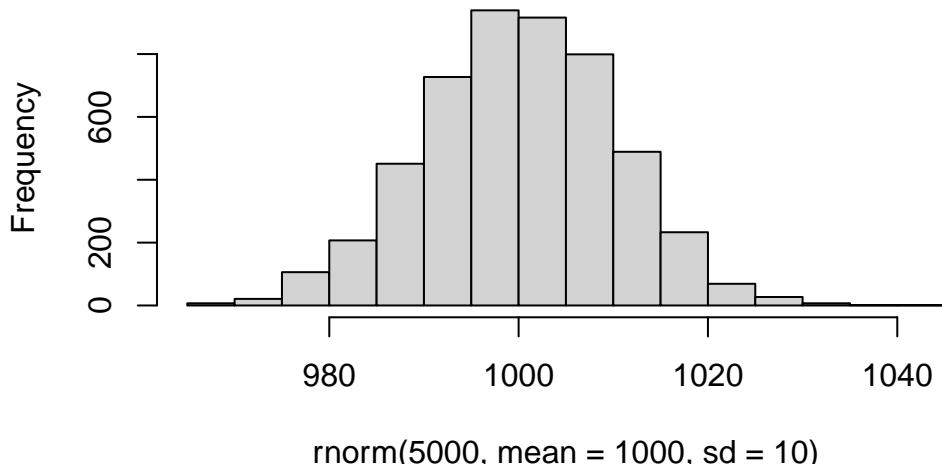
Today we will begin our exploration of some important machine learning methods, namely **clustering** and **dimensionality reduction**

Let's make up some input data where we know what the natural "clusters" are.

The function `rnorm()` can be useful here. ## K-means clustering

```
hist(rnorm(5000, mean=1000, sd= 10))
```

Histogram of rnorm(5000, mean = 1000, sd = 10)



Q. Generate 30 random numbers centered at +3 and another 30 at -3.

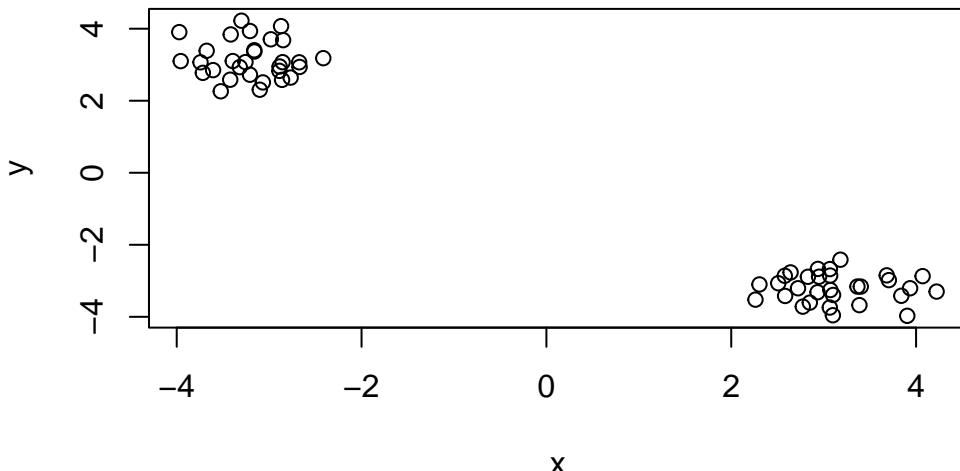
```
rnorm(30, mean=3, sd=0.5)
```

```
[1] 3.151270 2.878303 3.522507 3.294686 2.977663 2.484887 2.582619 3.007110  
[9] 3.180163 3.612721 3.690591 2.666465 3.142286 2.516513 2.630786 2.785373  
[17] 2.742216 3.469601 3.240985 3.169851 3.378167 2.650891 2.680185 3.005377  
[25] 2.841057 3.976550 3.660054 2.302960 3.398023 2.203407
```

```
rnorm(30, mean=-3, sd=0.5)
```

```
[1] -2.372569 -2.522981 -2.752637 -2.300521 -2.981157 -1.778568 -3.384532  
[8] -3.315427 -2.860868 -2.035190 -2.630400 -2.852770 -3.043116 -2.963215  
[15] -3.672775 -2.484820 -3.235167 -3.199280 -2.744795 -3.667816 -3.093506  
[22] -2.394515 -3.201059 -3.714107 -3.300802 -2.805318 -2.115911 -3.755147  
[29] -2.421337 -2.785730
```

```
tmp<- c(rnorm(30, mean=3, sd=0.5), rnorm(30, mean=-3, sd=0.5))  
  
x <- cbind(x=tmp, y=rev(tmp))  
plot(x)
```



K-means clustering

The main function in “base R” for k-means clustering is called surprisingly `kmeans()`:

```
km <- kmeans(x, 2)  
km
```

K-means clustering with 2 clusters of sizes 30, 30

Cluster means:

	x	y
1	3.135264	-3.197237
2	-3.197237	3.135264

Clustering vector:

Within cluster sum of squares by cluster:

```
[1] 12.39775 12.39775  
(between_SS / total_SS =  98.0 %)
```

Available components:

```
[1] "cluster"      "centers"       "totss"        "withinss"      "tot.withinss"  
[6] "betweenss"    "size"          "iter"         "ifault"
```

Q. What component of the results object details the cluster sizes?

km\$size

[1] 30 30

Q. What component of the results object details the cluster centers?

km\\$centers

```

          x      y
1  3.135264 -3.197237
2 -3.197237  3.135264

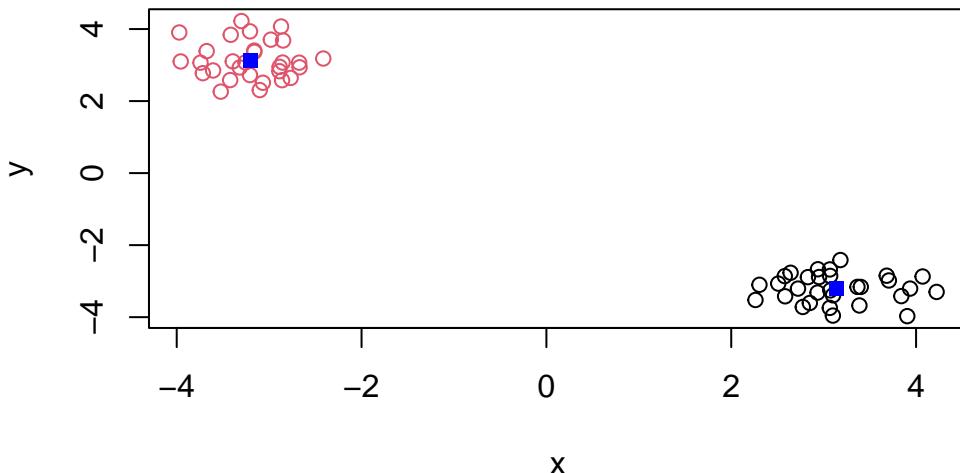
```

Q. What component of the results object details the cluster membership vector(i.e our main result of which points lie in which cluster)?

km\$cluster

Q. Plot our clustering results with points colored by cluster and also add the cluster centers as new points colored blue.

```
plot(x, col=km$cluster)
points(km$centers, col="blue", pch=15)
```



Q. `kmeans()` again but this time produce 4 clusters and call your resulting object `k4` and make a results figure like the one we made above.

```
k4 <- kmeans(x, 4)  
k4
```

```
K-means clustering with 4 clusters of sizes 10, 20, 20, 10
```

```
Cluster means:
```

	x	y
1	-3.259139	3.751668
2	-3.166286	2.827062
3	2.827062	-3.166286
4	3.751668	-3.259139

```
Clustering vector:
```

```
[1] 3 4 3 3 3 3 3 3 3 4 4 3 3 4 4 3 4 4 4 3 3 4 3 3 3 3 3 3 4 1 2 2 2 2 2 2 1  
[39] 2 2 1 1 1 2 1 1 2 2 1 1 2 2 2 2 2 2 2 2 1 2
```

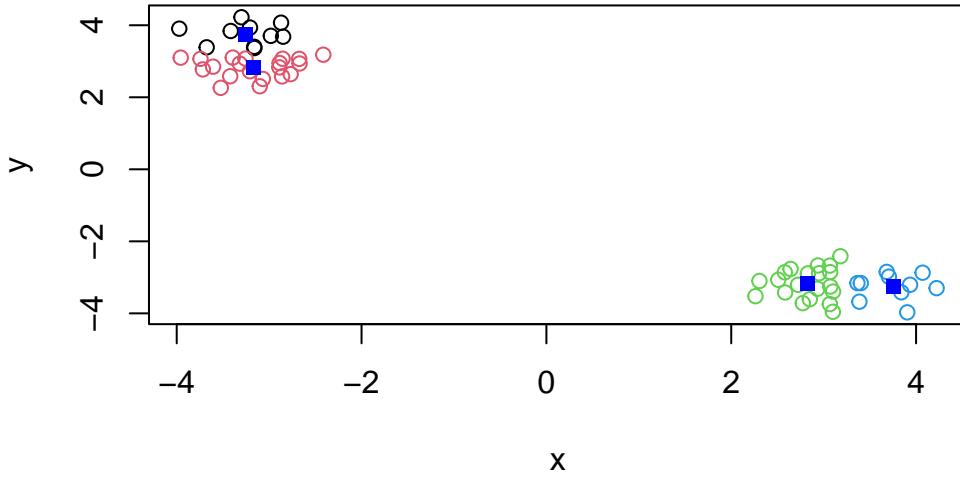
```
Within cluster sum of squares by cluster:
```

```
[1] 1.926699 4.714257 4.714257 1.926699  
(between_SS / total_SS = 98.9 %)
```

```
Available components:
```

```
[1] "cluster"      "centers"       "totss"        "withinss"      "tot.withinss"  
[6] "betweenss"    "size"         "iter"         "ifault"
```

```
plot(x, col=k4$cluster)  
points(k4$centers, col="blue", pch=15)
```



The Metric

```
km$tot.withinss
```

```
[1] 24.7955
```

```
k4$tot.withinss
```

```
[1] 13.28191
```

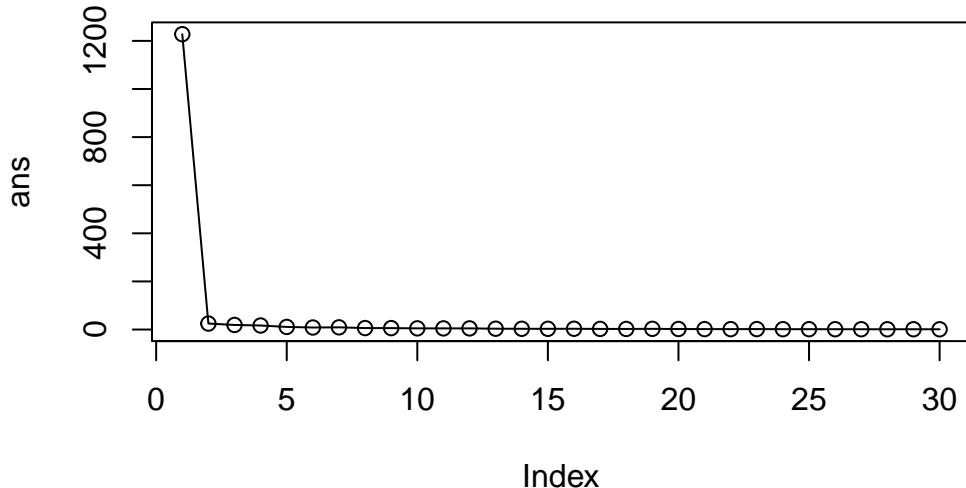
Q. Let's try different number of K (centers) from 1-30 and see what the best result is.

```
ans <- NULL
for(i in 1:30){
  ans <- c(ans, kmeans(x, centers=i)$tot.withinss)
}
ans
```

```
[1] 1227.8126552  24.7954977  19.0387051  16.6640108  10.8002881
[6]    8.4255939   9.1017182   6.0942003   5.8563116   4.8615038
[11]   4.6927577   4.6976621   3.4026164   3.1212455   2.9868517
```

```
[16] 3.2023638 2.5194383 2.4075950 2.8037573 2.2076444
[21] 1.7566220 1.6426979 1.8279377 1.3700086 1.3671189
[26] 1.2906851 1.2009566 1.0615792 1.1961547 0.9312107
```

```
plot(ans, typ="o")
```



`tot.withinss` shows how tight the clusters are. The lower the value the tighter the clusters.

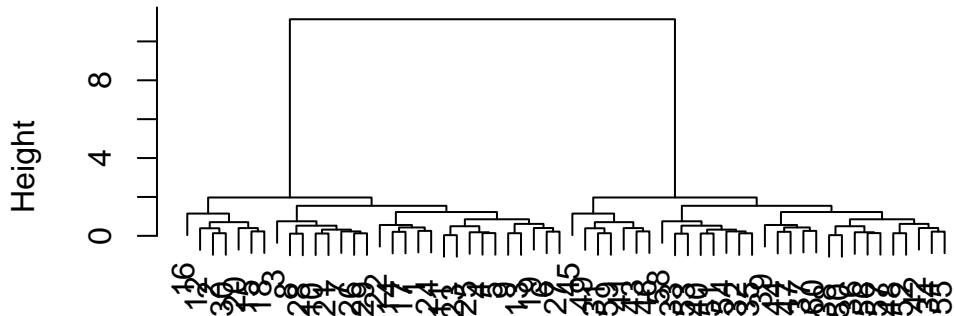
Key-point: ** K-means will impose a clustering structure on your data even if it is not there- it will always give the answer you asked for

Hierarchical CLustering

The main function for Hierarchical Clustering is called `hclust()`. Unlike `kmeans()` (which does all the work for us) you can't just pass `hclust()` our raw input data. It needs a “distance matrix” like the one returned from the `dist()` function.

```
d <- dist(x)
hc <- hclust(d)
plot(hc)
```

Cluster Dendrogram

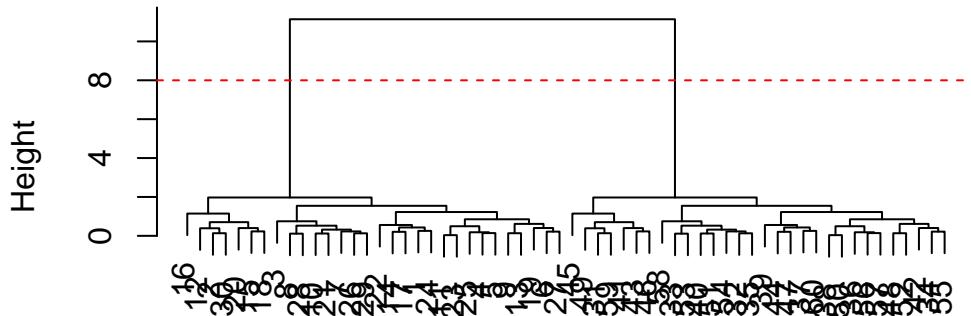


```
d  
hclust (*, "complete")
```

To extract our cluster membership vector from a `hclus` result object we have to “cut” our tree at a given height to yield separate “groups”/“branches”.

```
plot(hc)  
abline(h=8, col="red", lty=2)
```

Cluster Dendrogram



```
d  
hclust (*, "complete")
```

To do this we use the `cutree()` function on our `hclust()` object:

```
grps <- cutree(hc, h=8)  
grps
```

```
table(grps, km$cluster)
```

grps	1	2
1	30	0
2	0	30

PCA of UK food data

Import the data set of food consumption in the UK:

```
url <- "https://tinyurl.com/UK-foods"  
x <- read.csv(url)  
x
```

		X	England	Wales	Scotland	N.Ireland
1	Cheese		105	103	103	66
2	Carcass_meat		245	227	242	267
3	Other_meat		685	803	750	586
4	Fish		147	160	122	93
5	Fats_and_oils		193	235	184	209
6	Sugars		156	175	147	139
7	Fresh_potatoes		720	874	566	1033
8	Fresh_Veg		253	265	171	143
9	Other_Veg		488	570	418	355
10	Processed_potatoes		198	203	220	187
11	Processed_Veg		360	365	337	334
12	Fresh_fruit		1102	1137	957	674
13	Cereals		1472	1582	1462	1494
14	Beverages		57	73	53	47
15	Soft_drinks		1374	1256	1572	1506
16	Alcoholic_drinks		375	475	458	135
17	Confectionery		54	64	62	41

Q1. How many rows and columns are in your new data frame named x? What R functions could you use to answer this questions?

```
dim(x)
```

```
[1] 17 5
```

One solution to set the row names is to do it by hand.

```
rownames(x) <- x[,1]
rownames(x)
```

```
[1] "Cheese"           "Carcass_meat"    "Other_meat"
[4] "Fish"              "Fats_and_oils"   "Sugars"
[7] "Fresh_potatoes"   "Fresh_Veg"      "Other_Veg"
[10] "Processed_potatoes" "Processed_Veg" "Fresh_fruit"
[13] "Cereals"          "Beverages"     "Soft_drinks"
[16] "Alcoholic_drinks" "Confectionery"
```

```
x
```

		X	England	Wales	Scotland	N.Ireland
Cheese	Cheese	105	103	103	66	
Carcass_meat	Carcass_meat	245	227	242	267	
Other_meat	Other_meat	685	803	750	586	
Fish	Fish	147	160	122	93	
Fats_and_oils	Fats_and_oils	193	235	184	209	
Sugars	Sugars	156	175	147	139	
Fresh_potatoes	Fresh_potatoes	720	874	566	1033	
Fresh_Veg	Fresh_Veg	253	265	171	143	
Other_Veg	Other_Veg	488	570	418	355	
Processed_potatoes	Processed_potatoes	198	203	220	187	
Processed_Veg	Processed_Veg	360	365	337	334	
Fresh_fruit	Fresh_fruit	1102	1137	957	674	
Cereals	Cereals	1472	1582	1462	1494	
Beverages	Beverages	57	73	53	47	
Soft_drinks	Soft_drinks	1374	1256	1572	1506	
Alcoholic_drinks	Alcoholic_drinks	375	475	458	135	
Confectionery	Confectionery	54	64	62	41	

To remove the first column I can use the minus index trick.

```
x <- x[,-1]
x
```

	England	Wales	Scotland	N.Ireland
Cheese	105	103	103	66
Carcass_meat	245	227	242	267
Other_meat	685	803	750	586
Fish	147	160	122	93
Fats_and_oils	193	235	184	209
Sugars	156	175	147	139
Fresh_potatoes	720	874	566	1033
Fresh_Veg	253	265	171	143
Other_Veg	488	570	418	355
Processed_potatoes	198	203	220	187
Processed_Veg	360	365	337	334
Fresh_fruit	1102	1137	957	674
Cereals	1472	1582	1462	1494
Beverages	57	73	53	47
Soft_drinks	1374	1256	1572	1506
Alcoholic_drinks	375	475	458	135
Confectionery	54	64	62	41

A better way to do this is to set the row names to the first column by arguing with `read.csv()`.

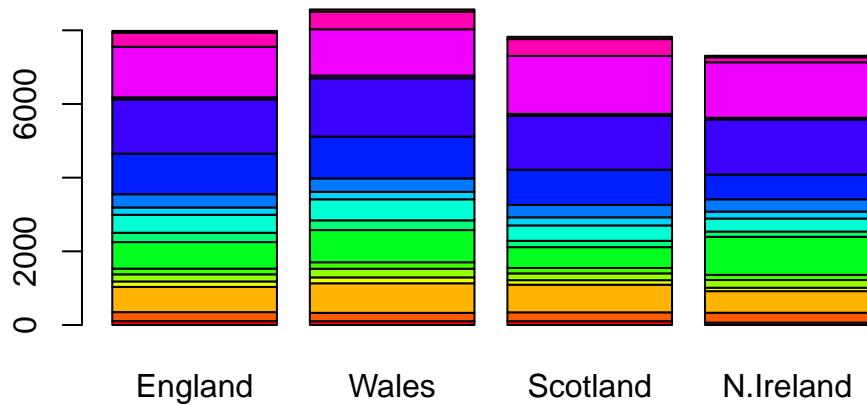
```
x <- read.csv(url, row.names=1)
x
```

	England	Wales	Scotland	N.Ireland
Cheese	105	103	103	66
Carcass_meat	245	227	242	267
Other_meat	685	803	750	586
Fish	147	160	122	93
Fats_and_oils	193	235	184	209
Sugars	156	175	147	139
Fresh_potatoes	720	874	566	1033
Fresh_Veg	253	265	171	143
Other_Veg	488	570	418	355
Processed_potatoes	198	203	220	187
Processed_Veg	360	365	337	334
Fresh_fruit	1102	1137	957	674
Cereals	1472	1582	1462	1494
Beverages	57	73	53	47
Soft_drinks	1374	1256	1572	1506
Alcoholic_drinks	375	475	458	135
Confectionery	54	64	62	41

Q2. Which approach to solving the ‘row-names problem’ mentioned above do you prefer and why? Is one approach more robust than another under certain circumstances?

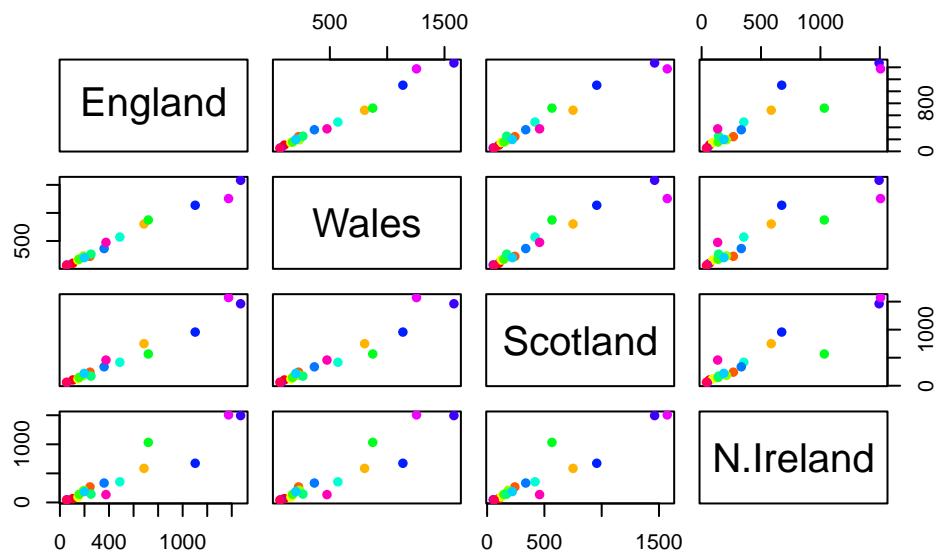
Spotting major differences and trends

```
barplot(as.matrix(x), beside=F, col=rainbow(nrow(x)))
```

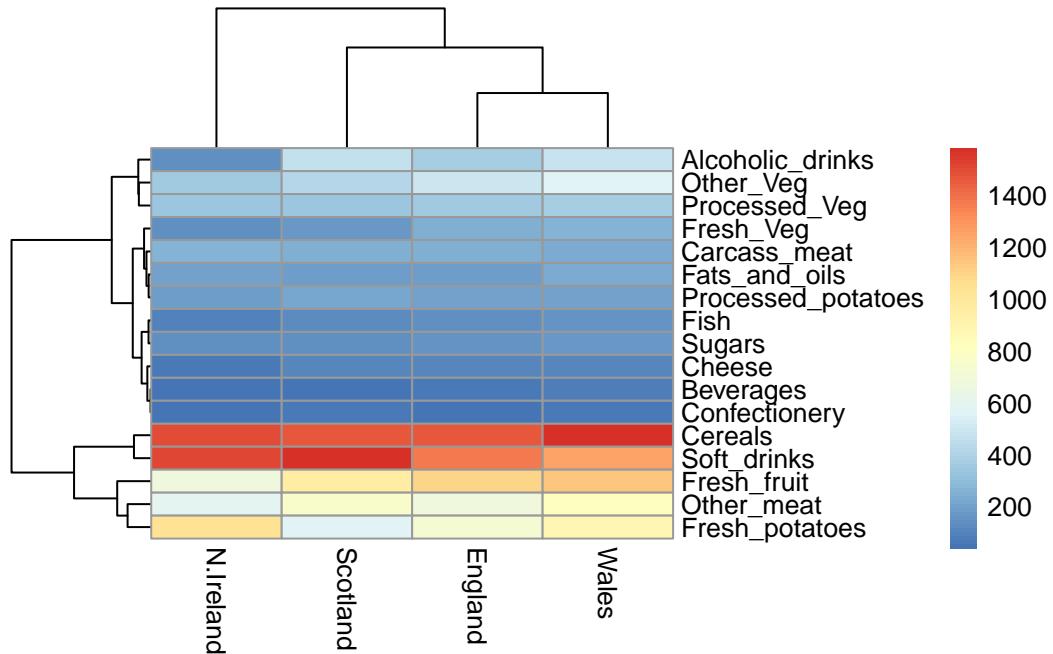


Pairs plots and heatmaps

```
pairs(x, col=rainbow(nrow(x)), pch=16)
```



```
library(pheatmap)
pheatmap( as.matrix(x) )
```



PCA to the rescue

The main PCA in “base R” is called `prcomp()`. This function wants the transpose of our food data as input(i.e the food as columns and the countries as rows).

```
pca <- prcomp(t(x))
pca
```

```
Standard deviations (1, ..., p=4):
[1] 3.241502e+02 2.127478e+02 7.387622e+01 3.175833e-14
```

```
Rotation (n x k) = (17 x 4):
```

	PC1	PC2	PC3	PC4
Cheese	-0.056955380	0.016012850	0.02394295	-0.694538519
Carcass_meat	0.047927628	0.013915823	0.06367111	0.489884628
Other_meat	-0.258916658	-0.015331138	-0.55384854	0.279023718
Fish	-0.084414983	-0.050754947	0.03906481	-0.008483145

```

Fats_and_oils      -0.005193623 -0.095388656 -0.12522257  0.076097502
Sugars            -0.037620983 -0.043021699 -0.03605745  0.034101334
Fresh_potatoes    0.401402060 -0.715017078 -0.20668248 -0.090972715
Fresh_Veg         -0.151849942 -0.144900268  0.21382237 -0.039901917
Other_Veg          -0.243593729 -0.225450923 -0.05332841  0.016719075
Processed_potatoes -0.026886233  0.042850761 -0.07364902  0.030125166
Processed_Veg      -0.036488269 -0.045451802  0.05289191 -0.013969507
Fresh_fruit        -0.632640898 -0.177740743  0.40012865  0.184072217
Cereals            -0.047702858 -0.212599678 -0.35884921  0.191926714
Beverages          -0.026187756 -0.030560542 -0.04135860  0.004831876
Soft_drinks         0.232244140  0.555124311 -0.16942648  0.103508492
Alcoholic_drinks   -0.463968168  0.113536523 -0.49858320 -0.316290619
Confectionery      -0.029650201  0.005949921 -0.05232164  0.001847469

```

```
summary(pca)
```

Importance of components:

	PC1	PC2	PC3	PC4
Standard deviation	324.1502	212.7478	73.87622	3.176e-14
Proportion of Variance	0.6744	0.2905	0.03503	0.000e+00
Cumulative Proportion	0.6744	0.9650	1.00000	1.000e+00

```
attributes(pca)
```

```

$names
[1] "sdev"      "rotation"   "center"     "scale"      "x"

$class
[1] "prcomp"

```

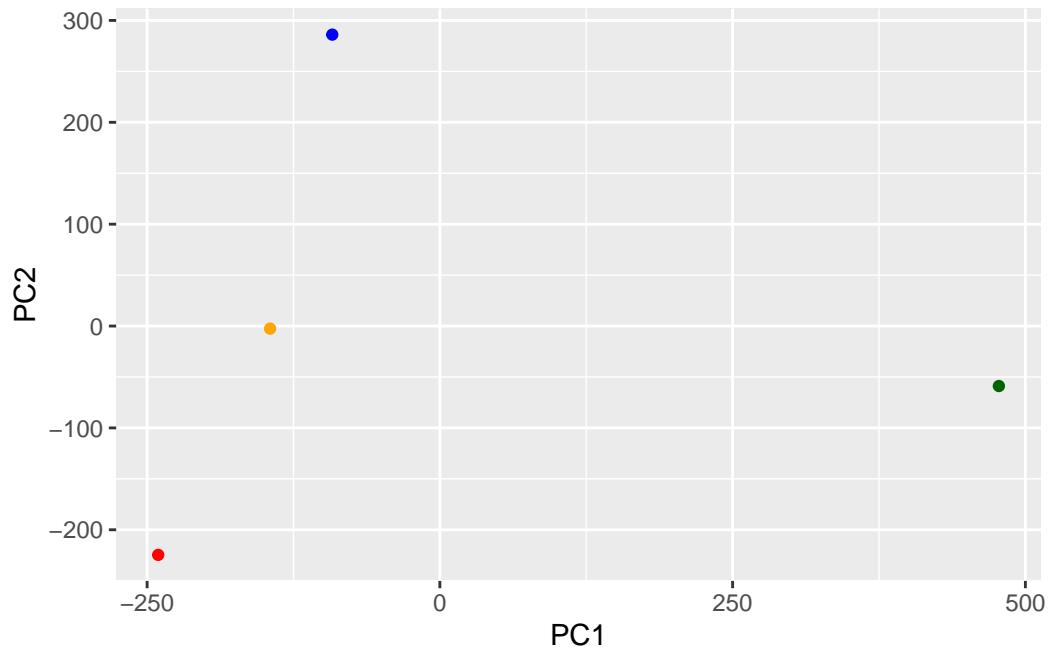
To make one of our main PCA result figures we turn to `pca$x` the scores along our new PCs. This is called “PC plot” or “score plot” or “Ordination plot” etc.

```
my_cols <- (c("orange","red","blue","darkgreen"))
```

```

library(ggplot2)
ggplot(pca$x) + aes(PC1,PC2)+geom_point(col=my_cols)

```



The second major result figure is called a “loadings plot” of “variable contributions plot” or “weight plot”.

```
ggplot(pca$rotation)+  
  aes(PC1, rownames(pca$rotation)) + geom_col()
```

